

Local Results Summary — Uniform-Sweep Template (0.9 / 0.8 / 0.7)

Blank fill-in template derived from local_results_summary.pdf. Table 1 is the consolidated “All values, row-by-row” reference with **all score cells intentionally empty** — one row per (model × benchmark), to be filled as the uniform $\alpha = 0.9 / 0.8 / 0.7$ evals land. Table 2 lists the four models that were previously blocked, updated to the current state (all now have on-disk data).

All values, row-by-row (consolidated reference) — blank template

One row per (model × benchmark). Score cells are intentionally empty: fill **Local** from each step-25 eval, **Paper** from BRV where reported, **Gap** = Local – Paper, and each Δ group = (+Uniform α) – Baseline. **n (Local) = n (BRV)** for every row (same MD5-matched TSVs); MathVerse-Overall 788* is the effective n (five modes share 788 problems).

					Baseline			+Uniform 0.9			+Uniform 0.8		Uniform 0.7		Δ 0.9 vs base		Δ 0.8	Δ 0.7
Model	Benchmark	n (Local)	n (BRV)	Type	Local	Paper	Gap	Local	Paper	Gap	Local	Local	Local	Paper	Local	Local		
LLaVA-Next-LLaMA3-8B (Dart-Prop)	MathVista-Overall	1000	1000	perception														
	MathVista-General (NEW)	460	460	perception														
	MathVista-Math (NEW)	540	540	reasoning														
	MathVerse-Overall	788*	788*	reasoning														
	MathVerse-T-D	788	788	reasoning														
	MathVerse-T-L	788	788	reasoning														
	MathVerse-V-I	788	788	reasoning														
	MathVerse-V-D	788	788	reasoning														
	MathVerse-V-O	788	788	reasoning														
	MathVision	304	304	reasoning														
Idefics2-8B (MAmmoTH-1)	MMStar-Overall	1500	1500	perception														
	MMStar-Math	250	250	reasoning														
	DynaMath-Average	5010	5010	reasoning														
	DynaMath-Worst (NEW)	501	501	reasoning														
	MathVista-Overall	1000	1000	perception														
	MathVista-General (NEW)	460	460	perception														
	MathVista-Math (NEW)	540	540	reasoning														
	MathVerse-Overall	788*	788*	reasoning														
	MathVerse-T-D	788	788	reasoning														
	MathVerse-T-L	788	788	reasoning														
	MathVerse-V-I	788	788	reasoning														
	MathVerse-V-D	788	788	reasoning														
	MathVerse-V-O	788	788	reasoning														
	MathVision	304	304	reasoning														
	MMStar-Overall	1500	1500	perception														
	MMStar-Math	250	250	reasoning														
	DynaMath-Average	5010	5010	reasoning														
	DynaMath-Worst (NEW)	501	501	reasoning														
	Qwen2-VL-7B-Instruct (Qwen2-Math-7B)	MathVista-Overall	1000	1000	perception													

[illegible]

[illegible]

[illegible]

Previously-blocked models — now resolved

All four models that completed the 9-model enrichment-target set previously had no step-25 score CSVs. As of this update **every one now has on-disk eval data**, so their blockers are resolved. CSV counts are score/acc files under results/25-merge/<model>/ at update time.

Model	Auto-routed math LLM	Prior blocker	Current status
LLaVA-OneVision-7B	Qwen2-Math-7B (HF)	VLMEvalKit_brv venv missing the <i>llava</i> Python package; eval jobs died with No module named 'llava' (also hit LLaVA-1.5-7B).	Resolved. 24 merges, 402 score/acc CSVs on disk. (Handoff: notes/vlmevalkit_handoff_llava_next_install.md.)
Qwen2.5-VL-7B	Qwen2.5-Math-7B (HF)	(1) hf-xet Rust panic mid-download of Qwen2.5-Math-7B; (2) eval KeyError: 'Qwen2.5-VL-7B-Instruct' (missing from VLMEvalKit supported_VLM).	Resolved. 24 merges, 357 score/acc CSVs on disk. (Handoff: notes/vlmevalkit_handoff_qwen25vl_7b_registry.md.)
InternVL2.5-8B	internlm2-math-plus-7b (HF)	Merge once killed by TERM_THREADLIMIT + TypeError: dtype (both fixed); eval failed with 'InternLM2ForCausalLM' object has no attribute 'generate' (GenerationMixin not inherited).	Resolved. 24 merges, 449 score/acc CSVs on disk. (Handoff: notes/internvl_handoff_generation_mixin.md.)
Qwen2.5-VL-3B	Qwen2.5-3B-DAPO-math-reasoning	Hard-blocked (run_pipeline.sh exit 1): official Qwen2.5-Math ships only 1.5B/7B (28L), shape-incompatible with the 36-layer Qwen2.5-3B backbone, so no math donor was wired in.	Resolved. Donor jaygala24/Qwen2.5-3B-DAPO-math-reasoning (36-layer, shape-compatible) added to step 25; 4 merges, 74 score/acc CSVs on disk.